

SCP wiki data analysis: Average ratings of authors

```
# Data for all SCP articles
data1 <- read.table("data-3-09-21.txt", header=T)

# Data for author A
data2 <- read.table("author-data.txt", header=T)
```

1 Summary

Author A is one of my favourite authors on the site. The goal of this analysis is to determine whether their entries are more well-received by the general SCP wiki community than other authors' entries.

The average rating of entries belonging to author A (average 184.9) is greater than the average rating of entries belonging to other authors (average 168.83569). The difference is probably not significant, though these results should be interpreted skeptically, due to a number of reasons that led to the model used being very unsuitable for the analysis:

- Negatively-rated entries are almost always removed from the site, causing the data to be skewed rightward.
- The skill level of different authors and the quality of their entries may vary.
- The rating of an entry depends on its visibility and quality. Some popular SCP entries may reference other entries, causing the visibility of one to depend on the visibility of another. As well, more skilled authors will tend to produce higher-quality entries.

The average rating-comment ratio of entries belonging to author A (average 6.9191865) is greater than the average rating-comment ratio of other authors (average 4.862698). The difference is significant. The model used for the analysis was somewhat unsuitable for the analysis, due to the right-skewness of the data caused by negatively-rated entries being removed from the site. Therefore these results should also be interpreted with some skepticism.

2 Calculations

2.1 Preprocessing

Assign entries not belonging to author A to group 1, and entries belonging to author A to group 2. SCP-001 has been removed from the data, since it has multiple proposals, and the rating and comments recorded are for the hub page for the proposals, and do not correspond to any particular author's entry. To prevent division by 0, 1 was added to the denominator in the calculation of the rating-comment ratios.

```
head(all.data)
```

```
##   scp rating comments group      ratio
## 2    2    1527      116     1 13.051282
## 3    3     678       86     1  7.793103
## 4    4     991      122     1  8.056911
## 5    5     566      102     1  5.495146
## 6    6     508      100     1  5.029703
## 7    7     496       47     1 10.333333
```

```

head(group1)

##    scp rating comments group      ratio
## 2    2    1527      116     1 13.051282
## 3    3     678       86     1  7.793103
## 4    4     991      122     1  8.056911
## 5    5     566      102     1  5.495146
## 6    6     508      100     1  5.029703
## 7    7     496       47     1 10.333333

head(group2)

##    scp rating comments group      ratio
## 670 670     294       35     2 8.166667
## 737 737     167       32     2 5.060606
## 753 753     302       41     2 7.190476
## 777 777     125       22     2 5.434783
## 779 779     160       33     2 4.705882
## 844 844      65       16     2 3.823529

nrow(all.data) # Total number of observations

## [1] 5697

nrow(group1) # Number of replicates in group 1

## [1] 5587

nrow(group2) # Number of replicates in group 2

## [1] 110

mean(group1$rating)

## [1] 168.8357

mean(group2$rating)

## [1] 184.9

mean(group1$rating / (group1$comments + 1))

## [1] 4.862698

mean(group2$rating / (group2$comments + 1))

## [1] 6.919186

```

2.2 Testing whether the the difference in ratings is significant

The average rating of entries belonging to author A is greater than the average rating of entries not belonging to author A. This section measures the significance of the difference, by estimating the attributes of a hypothetical random process that produces the rating for each entry.

2.2.1 Model

The separation of the entries into two groups and an unequal number of entries in each of these groups would suggest an unbalanced completely randomized design:

$$Y_{ij} = \mu + \tau_i + R_{ij} \quad R_{ij} \sim N(0, \sigma^2)$$

where

- Y_{ij} is the response variable, corresponding to the random process for the rating of an entry
- $i = 1, 2$ correspond to the treatment groups (entries not belonging to author A, and entries belonging to author A).
- $j = 1, 2, \dots, 5587$ are the replicates for group 1 (entries not belonging to author A), and $j = 1, 2, \dots, 110$ are the replicates for group 2 (entries belonging to author A)
- μ is the mean rating
- τ_1 and τ_2 are the treatment effects corresponding to group 1 (entries not belonging to author A) and group 2 (entries belonging to author A) respectively.

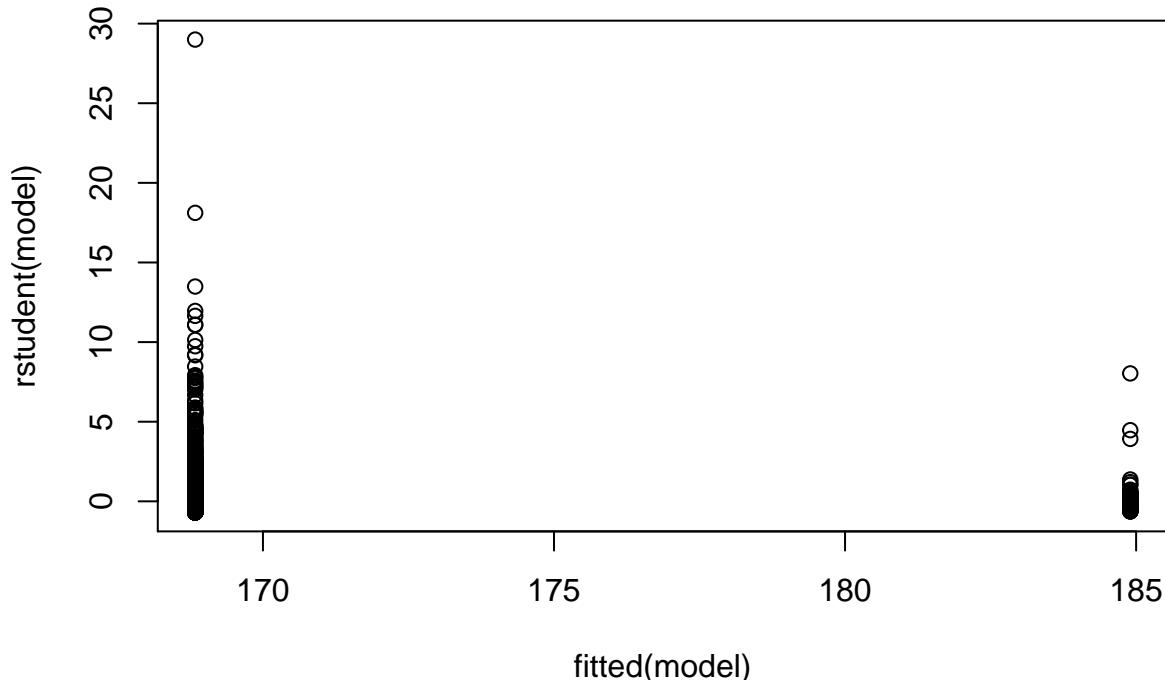
There is one constraint on the model: $0 = 5587\tau_1 + 110\tau_2$.

2.2.2 Assessing appropriateness of model

```
model <- lm(rating~group, all.data)
sigmahat <- summary(model)$sigma
sigmahat # Residual standard error
```

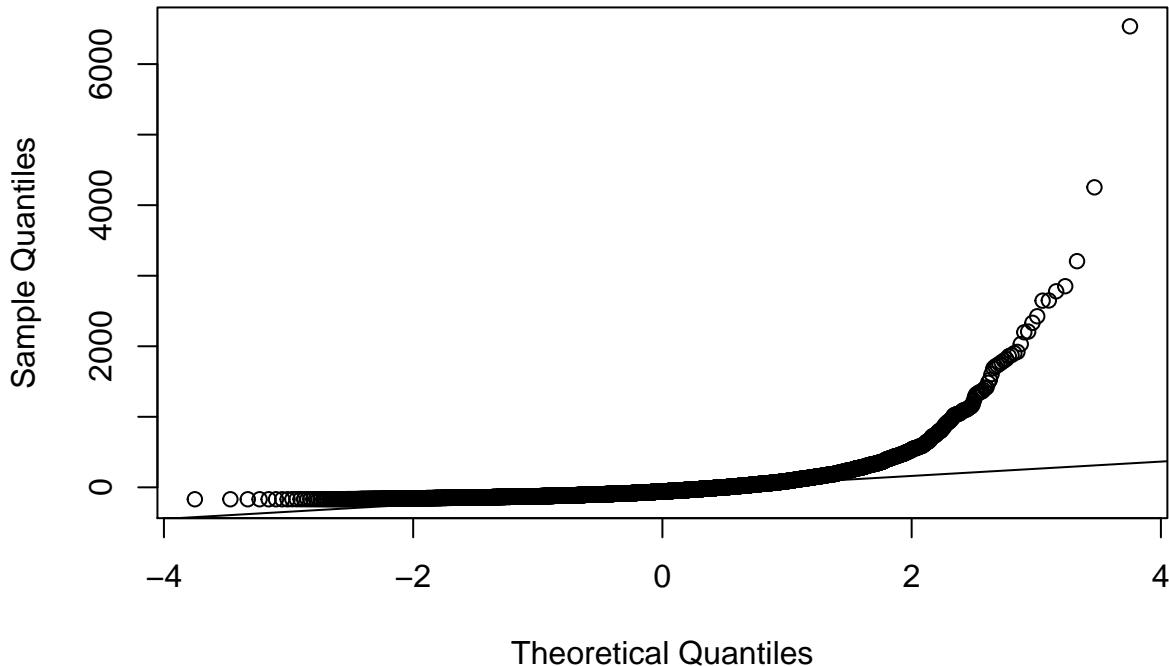
```
## [1] 241.438
```

```
plot(fitted(model), rstudent(model))
```



```
qqnorm(resid(model))
qqline(resid(model))
```

Normal Q–Q Plot



This is not an appropriate model. The assumptions of a linear model do not appear to be satisfied:

- The data is right-skewed, indicating that the distribution is probably not Normal. This is likely because negatively-rated entries are almost always removed from the site.
- The variance in the residuals for author A's entries appears to be less than the variance in the residuals for other authors' entries, so the constant variance assumption might not hold. A higher variance in the other authors' entries may be due to varying levels of skill among the other authors. This could be remedied by adding factor levels for other authors. However, the consistency of the quality of an author's entries likely differs between different authors, leading to different variances in the ratings of entries belonging to each author, so even if this were done, the constant variance assumption might still not hold.
- The ratings of each entry also may not be independent. The rating of an entry depends on its visibility and quality. Some popular SCP entries may reference other entries, causing the visibility of one to depend on the visibility of another. As well, more skilled authors will tend to produce higher-quality entries, which may be another source of dependence among the data.

2.2.3 Hypothesis test

If we were to perform a hypothesis test anyway, we would be interested in the attribute $\theta = \tau_2 - \tau_1$, and the hypothesis $H_0 : \theta \leq 0$ vs. $H_a : \theta > 0$. The estimator is $\tilde{\theta} = \tilde{\tau}_2 - \tilde{\tau}_1 = \frac{1}{r_2} \sum_{j=1}^{r_2} Y_{2j} - \frac{1}{r_1} \sum_{j=1}^{r_1} Y_{1j}$.

$$\begin{aligned}
E(\tilde{\theta}) &= E \left(\tilde{\theta} = \tilde{\tau}_2 - \tilde{\tau}_1 = \frac{1}{r_2} \sum_{j=1}^{r_2} Y_{2j} - \frac{1}{r_1} \sum_{j=1}^{r_1} Y_{1j} \right) \\
&= \frac{1}{r_2} \sum_{j=1}^{r_2} E(Y_{2j}) - \frac{1}{r_1} \sum_{j=1}^{r_1} E(Y_{1j}) \\
&= \frac{1}{r_2} \sum_{j=1}^{r_2} (\mu + \tau_2) - \frac{1}{r_1} \sum_{j=1}^{r_1} (\mu + \tau_1) \\
&= \mu + \tau_2 - \mu - \tau_1 \\
&= \theta
\end{aligned}$$

$$\begin{aligned}
Var(\tilde{\theta}) &= Var \left(\frac{1}{r_2} \sum_{j=1}^{r_2} Y_{2j} - \frac{1}{r_1} \sum_{j=1}^{r_1} Y_{1j} \right) \\
&= \frac{1}{r_2^2} \sum_{j=1}^{r_2} Var(Y_{2j}) + \frac{1}{r_1^2} \sum_{j=1}^{r_1} Var(Y_{1j}) \\
&= \frac{1}{5587} \sigma^2 + \frac{1}{110} \sigma^2
\end{aligned}$$

$\tilde{\theta}$ is a linear combination of Normal random variables, so it's Normally distributed. Therefore, assuming $\theta = 0$,

$$\begin{aligned}
\tilde{\theta} &\sim N \left(0, \left(\frac{1}{5587} + \frac{1}{110} \right) \sigma^2 \right) \\
\frac{\tilde{\theta}}{\left(\frac{1}{5587} + \frac{1}{110} \right) \sigma^2} &\sim N(0, 1) \\
\frac{\tilde{\theta}}{\left(\frac{1}{5587} + \frac{1}{110} \right) \tilde{\sigma}^2} &\sim t_{n-3+1}
\end{aligned}$$

where $n = 5697$ is the number of observations; there are 3 non- σ parameters and 1 constraint. The pivotal quantity is

$$D = \frac{\tilde{\theta}}{\left(\frac{1}{5587} + \frac{1}{110} \right) \tilde{\sigma}^2} \sim t_{5695}$$

The discrepancy statistic is

$$d = \frac{\hat{\theta}}{\left(\frac{1}{5587} + \frac{1}{110} \right) \hat{\sigma}^2}$$

The p-value would be $p = Pr(D > d)$.

```

thetahat <- mean(group2$rating) - mean(group1$rating)
d <- thetahat / ((1 / 5587 + 1 / 110) * sigmahat ^ 2)
pvalue <- 1 - pt(d, 5695)
pvalue

## [1] 0.4881422

```

The p-value 0.4881422 is greater than 0.1, so there is insufficient evidence against H_0 . The difference between the average rating of author A's entries and the average rating of other entries is not statistically significant. Since the model is not an appropriate one, this result should be interpreted with skepticism.

2.3 Testing whether the difference in rating-comment ratio is significant

The average rating-comment ratio of author A is greater than the average rating-comment ratio of other authors on the site. This section measures the significance of the difference, by estimating the attributes of a hypothetical random process that produces the ratio for each entry.

2.3.1 Model

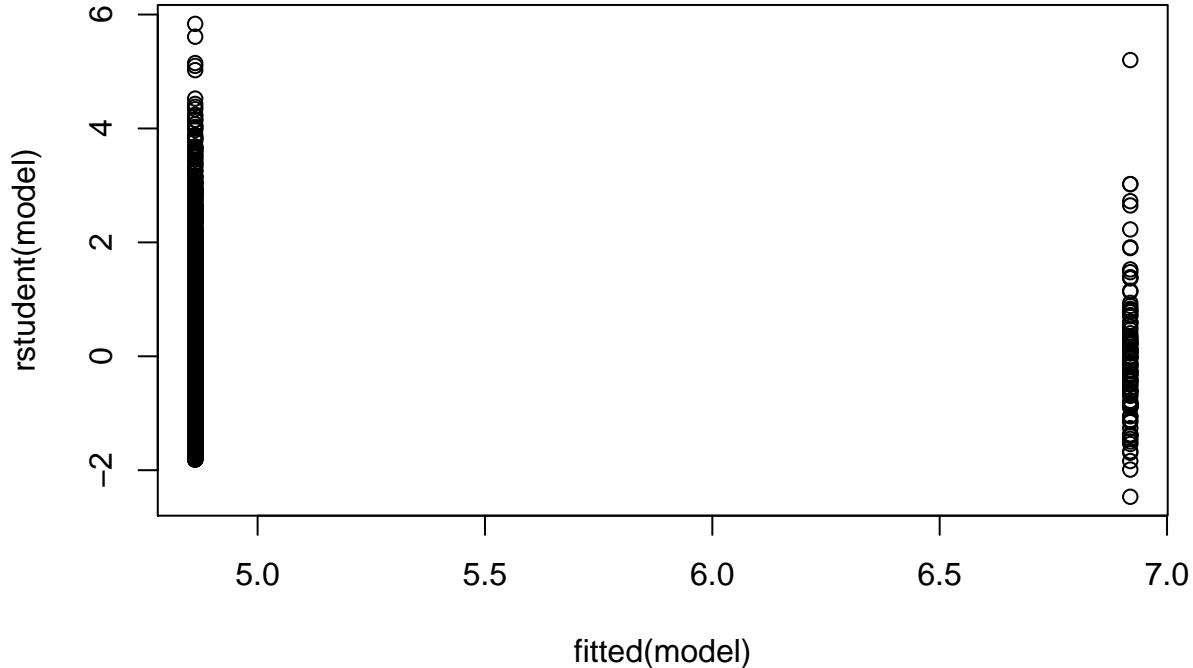
The model will be the same one as before, aside from a difference in interpretation of the attribute μ : it is now the mean rating-comment ratio.

2.3.2 Assessing appropriateness of model

```
model <- lm(ratio~group, all.data)
sigmahat <- summary(model)$sigma
sigmahat # Residual standard error

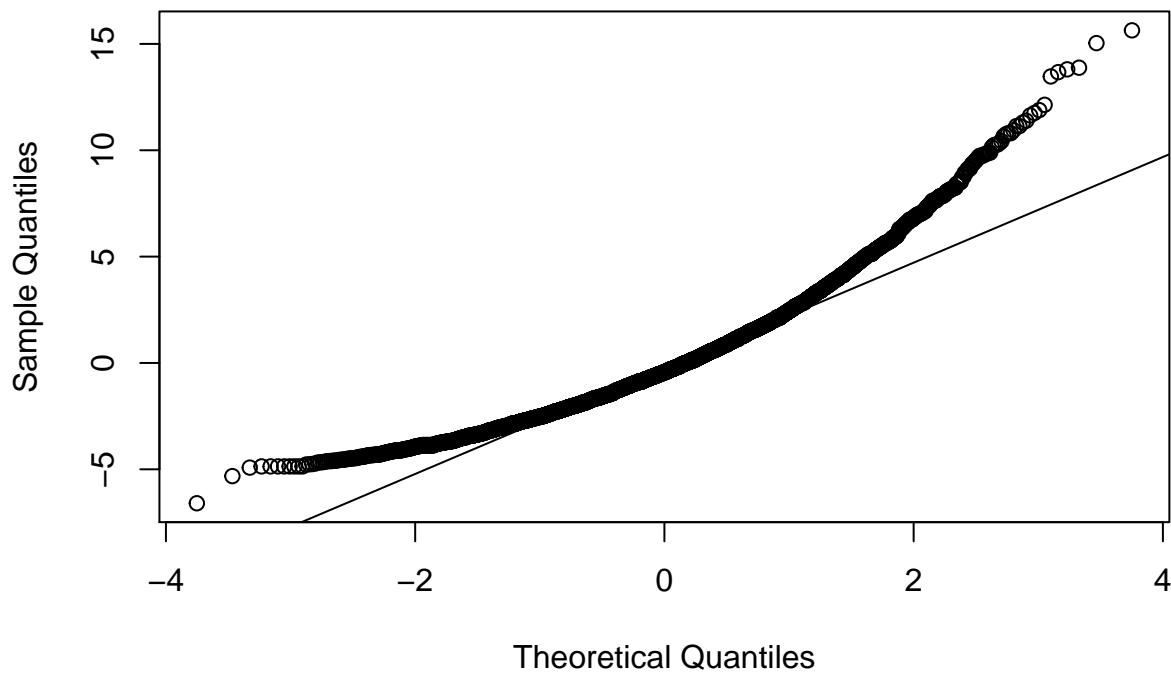
## [1] 2.687509

plot(fitted(model), rstudent(model))
```



```
qqnorm(resid(model))
qqline(resid(model))
```

Normal Q-Q Plot



The model is not suitable, mainly due to the right-skewness of the data, which indicates that the distribution of the residuals is probably not Normal. The plot of the residuals indicates that the expectation assumption ($E(R_{ij}) = 0$ for all i, j) also might not hold.

2.3.3 Hypothesis test

If we were to perform a hypothesis test anyway, we would be interested in the attribute $\theta = \tau_2 - \tau_1$, and the hypothesis $H_0 : \theta \leq 0$ vs. $H_a : \theta > 0$. The math is the same as before.

```
thetahat <- mean(group2$ratio) - mean(group1$ratio)
d <- thetahat / ((1 / 5587 + 1 / 110) * sigmahat ^ 2)
pvalue <- 1 - pt(d, 5695)
pvalue
```

```
## [1] 0
```

The p-value 0 is less than 0.01 (it is probably too small for float precision to properly display its value). There is strong evidence against H_0 . The difference between the rating-comment ratios of author A and other authors is significant. Due to the model not being appropriate, these results should be interpreted with skepticism.